

Deriving a Kronecker-Free Functional Quantum Simulator

Martin Elsman

Department of Computer Science,
University of Copenhagen (DIKU)

WQS 2025

Seoul, Republic of Korea


June 17, 2025



Paper pdf

Quantum Circuit Semantics and State-Vector Simulators

- The semantics of quantum circuits are specified **declaratively** and **compositionally** as unitary matrices.
- Efficient quantum simulators are **imperative** and operate using in-place array-updates on a global state vector with indices calculated using complex addressing.
- A formal connection is seldom established!



Contributions:

- We derive, from the semantics of circuits, an interpreter that works directly on state vectors without first generating large unitary matrices.
- The derivation builds on the well-known concept of *vectorisation* (i.e., a connection between column stacking and the Kronecker product).

Qubits, Gates, and Multi-Qubit States

- A *qubit* can be modelled as a two-dimensional complex vector $\begin{bmatrix} \alpha \\ \beta \end{bmatrix}$ specifying a linear combination $\alpha|0\rangle + \beta|1\rangle$ of the basis vectors $|0\rangle$ and $|1\rangle$ such that $|\alpha|^2 + |\beta|^2 = 1$.
- *Single-qubit gates* model transformations of a qubit and can be represented as 2×2 unitary complex matrices U , meaning $U^\dagger U = UU^\dagger = I$ (norm-preserving and reversible):

$$X = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix} \quad Y = \begin{bmatrix} 0 & -i \\ i & 0 \end{bmatrix} \quad Z = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix} \quad H = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix} \quad T = \begin{bmatrix} 1 & 0 \\ 0 & e^{i\pi/4} \end{bmatrix} \quad I = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$$

- A state of more qubits is modelled as a product of all standard bases.
 - E.g., a two-qubit state is $\alpha|00\rangle + \beta|01\rangle + \gamma|10\rangle + \theta|11\rangle$, where $|\alpha|^2 + |\beta|^2 + |\gamma|^2 + |\theta|^2 = 1$. I.e., a four-element complex vector $[\alpha \ \beta \ \gamma \ \theta]^T$.
 - The state space doubles with every added qubit!

Multi-Qubit Circuits (Composition of Gates)

Multi-qubit circuits are modelled using **matrix multiplication** (horizontal composition) and **tensor products** (vertical composition).

$$A \otimes B = \begin{bmatrix} a_{11}B & \cdots & a_{1n}B \\ \vdots & \ddots & \vdots \\ a_{m1}B & \cdots & a_{mn}B \end{bmatrix} \quad A \otimes \mathbf{1} = A = \mathbf{1} \otimes A$$

Special matrices are used for specifying **qubit swapping** (SW) and for specifying so-called **control gates** ($C_n U$).


$$SW = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad I_n = \begin{cases} \mathbf{1} & \text{if } n = 0 \\ I \otimes I_{n-1} & \text{otherwise} \end{cases} \quad C_0 U = U$$
$$C_n U = \begin{bmatrix} I_n & 0 \\ 0 & C_{n-1} U \end{bmatrix}$$

A Language for Specifying Circuits

We can define a small DSL for specifying circuits:

```
c ::= X | Y | Z | H | T | I | SW  
| c + c | c ⊗ c | C c
```

Circuits may be composed sequentially (using $+$) and vertically (using \otimes).



(Example circuit)

```
val c0 = I ⊗ H + C X + Z ⊗ Z + C X + I ⊗ H  
val c1 = c0 ⊗ I + I ⊗ SW + C X ⊗ Y
```

(Circuit code)

Semantics (denotation)

We can define a partial function $h : Circ \rightarrow \mathbb{N}$ for computing the height of a circuit c :

$$\begin{aligned} h(c_1 \otimes c_2) &= h c_1 + h c_2 & h(c_1 + c_2) &= h c_1 & (h c_1 = h c_2) \\ h(C c) &= 1 + h c & h c &= 1 & (c \in \{X, Y, Z, H, T, I\}) \end{aligned}$$

The semantics of a circuit c with height $k = h c$ is specified by the unitary matrix $\llbracket c \rrbracket : \mathbb{C}^{2^k \times 2^k}$ defined as follows:

$$\begin{aligned} \llbracket c_1 + c_2 \rrbracket &= \llbracket c_2 \rrbracket \llbracket c_1 \rrbracket & \llbracket X \rrbracket &= X & \llbracket H \rrbracket &= H \\ \llbracket c_1 \otimes c_2 \rrbracket &= \llbracket c_2 \rrbracket \otimes \llbracket c_1 \rrbracket & \llbracket Y \rrbracket &= Y & \llbracket T \rrbracket &= T \\ \llbracket C c \rrbracket &= C_{(h c)}(\llbracket c \rrbracket) & \llbracket Z \rrbracket &= Z & \llbracket I \rrbracket &= I \end{aligned}$$

For evaluating a circuit c on a state vector $v : \mathbb{C}^{2^{(h c)}}, we can compute $\llbracket c \rrbracket v$.$

Can we do better?

Kronecker-Free Interpretation (I)

We can *derive* an efficient interpreter $\mathcal{I}[c] v$ inductively:

For any circuit c such that $h c = k$ and state vector $v : \mathbb{C}^{2^k}$, we have $\mathcal{I}[c] v = \llbracket c \rrbracket v$

The derivation is based on *vectorisation* and properties of matrix multiplication:

$$(A \otimes B) v = \text{vec}(B(\text{unvec } v)A^T) \quad (\text{vectorisation})$$

$$\text{vec } A = \text{flatten } (A^T) \quad (\text{column stacking})$$

$$\text{unvec } v = (\text{unflatten } v)^T \quad (\text{cut vector into columns})$$

$$AB = (\text{map } (\lambda v. A v) B^T)^T \quad (\text{matmul functionalise})$$

$$(AB)^T = B^T A^T \quad (\text{matmul transpose})$$

Kronecker-Free Interpretation (II)

We proceed by case analysis.

For base cases, the denotation gives us the desired implementations immediately:

$$\mathcal{I}[c] v = v$$

$$c = \mathbf{I}$$

$$\mathcal{I}[c] v = \|c\| v$$

$$c \in \{\mathbf{X}, \mathbf{Y}, \mathbf{Z}, \mathbf{H}, \mathbf{T}, \mathbf{SW}\}$$

For sequential composition, we have

$$\begin{aligned} \mathcal{I}[c_1 + c_2] v &= \|c_1 + c_2\| v && \text{from property} \\ &= (\|c_2\| \|c_1\|) v && \text{from denotation} \\ &= \|c_2\| (\|c_1\| v) && \text{matrix-vector multiply} \\ &= \|c_2\| (\mathcal{I}[c_1] v) && \text{by induction} \\ &= \mathcal{I}[c_2] (\mathcal{I}[c_1] v) && \text{by induction} \end{aligned}$$

Kronecker-Free Interpretation (III)

For the tensor-product case, we have

$$\begin{aligned}\mathcal{I}[c_1 \otimes c_2] v &= [\![c_1 \otimes c_2]\!] v && \text{from property} \\&= ([\![c_1]\!] \otimes [\![c_2]\!]) v && \text{from denotation} \\&= \text{vec}([\![c_2]\!])(\text{unvec } v)([\![c_1]\!]^T) && \text{vectorisation} \\&= \text{vec}(([[\![c_1]\!]]([\![c_2]\!])(\text{unvec } v))^T)^T && \text{matmul transpose} \\&= \text{let } W = \text{map } [\![c_2]\!] (\text{unvec } v)^T && \text{matmul functionalise twice} \\&\quad \text{let } Y = \text{map } [\![c_1]\!] W^T \\&\quad \text{in vec } Y \\&= \text{let } W = \text{map } (\mathcal{I}[c_2]) (\text{unvec } v)^T && \text{by induction twice} \\&\quad \text{let } Y = \text{map } (\mathcal{I}[c_1]) W^T \\&\quad \text{in vec } Y\end{aligned}$$

Full Kronecker-Free Interpreter in a Functional Language


```
fun eval (c:c, v:vec) : vec =
  M.matvecmul_gen C.* C.+ (C.fromInt 0) (sem c) v

fun interp (c:c) (v:vec) : vec =
  case c of I => v
  | A+B => interp B (interp A v)
  | A⊗B => let val V = unvec (pow2(height A),pow2(height B)) v
              val W = mapRows (interp B) (M.transpose V)
              val Y = mapRows (interp A) (M.transpose W)
              in vec Y
             end
  | C A => Vector.concat [vecTake (pow2(height A)) v,
                           interp A (vecDrop (pow2(height A)) v)]
  | _ => eval (c,v)
```

Recall: A controlled circuit $C A$ acts as the identity on the first half of a state vector and as A on the second half.

Benchmarks

We have implemented Grover's search algorithm and measured the running time for different variations of simulators and different sizes of search problems.¹



¹Compiled with MLton and executed on a 2023 MacBook Pro (M2 Max) with 32GB of memory.

Much Related Work

- Many other simulators (qsim, QuEST, ...)
- Deriving a data-parallel gate-library (ARRAY '25: Gate Fusion is Map Fusion).

Opportunities and Future Work

- Performance comparisons with state-of-the-art simulators.
- Monadic generation of circuit-specialised simulators (e.g., for generating GPU code).
- Specialised code for diagonal and antidiagonal circuits.
- Offline gate fusion.
- More basic gates, measurements, and support for swapping distant qubits.
- ...