

# Gate Fusion is Map Fusion

Martin Elsman<sup>1</sup>

Department of Computer Science,  
University of Copenhagen (DIKU)

ARRAY 2025

Seoul, Republic of Korea

June 17, 2025



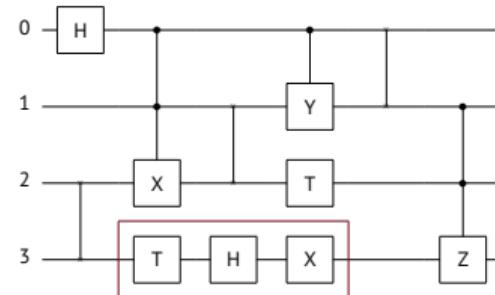
Paper pdf

---

<sup>1</sup>Joined work with Troels Henriksen, DIKU

## State-Vector Quantum Simulators and Gate Fusion

- Efficient quantum circuit simulators are **imperative**.
- Gates operate on a **global state vector** (of size  $2^k$ ) with the precise access and update patterns depending on which qubits the gates operate on.
- Gate-fusion is **important for performance**.



(Example circuit,  $k = 4$ )

**Contribution:** Using the well-known concept of *vectorisation* (i.e., column stacking meets the Kronecker product), we demonstrate that *gate operations become index-transformed maps* and *gate fusion becomes map fusion*.

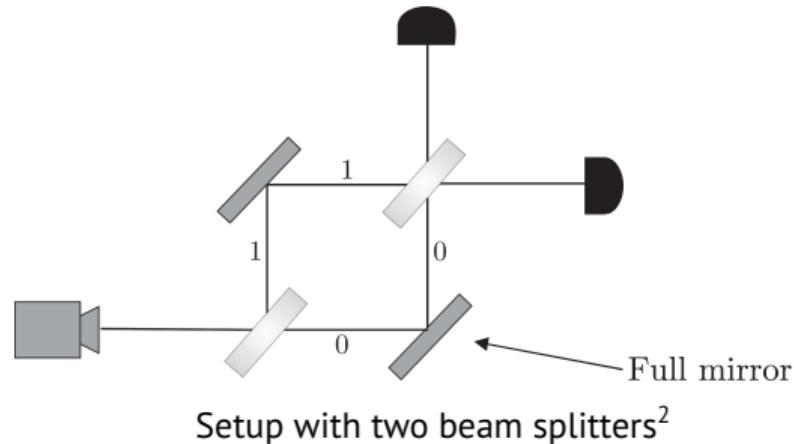
Efficient quantum simulators set the bar for quantum supremacy and are used for exploring quantum algorithms (qsim, QuEST, ...).

## The Breakdown of Classical Reasoning: The Double Beam Splitter

One beam splitter puts a photon in “superposition”, but two beam splitters in sequence act as the identity!

A beam splitter can be modelled by a so-called Hadamard gate ( $H$ ):

$$\begin{aligned} HH &= \left( \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix} \right) \left( \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix} \right) \\ &= \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix} \\ &= I \end{aligned}$$



<sup>2</sup>From Phillip Kaye, Raymond Laflamme, and Michele Mosca. *An Introduction to Quantum Computing*. Oxford University Press. 2007

## Qubits and Single-Qubit Gates

- A *qubit* can be modelled as a two-dimensional complex vector  $\begin{bmatrix} \alpha \\ \beta \end{bmatrix}$  specifying a linear combination  $\alpha|0\rangle + \beta|1\rangle$  of the basis vectors  $|0\rangle$  and  $|1\rangle$  such that  $|\alpha|^2 + |\beta|^2 = 1$ .
- A *single-qubit gate* models a transformation of a qubit and can be represented as a  $2 \times 2$  unitary complex matrix  $U$ , meaning  $U^\dagger U = UU^\dagger = I$  (norm-preserving and reversible).
- Pauli-gates:

$$X = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix} \quad Y = \begin{bmatrix} 0 & -i \\ i & 0 \end{bmatrix} \quad Z = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$$

- Hadamard-gate and T-gate:

$$H = \frac{1}{\sqrt{2}} \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix} \quad T = \begin{bmatrix} 1 & 0 \\ 0 & e^{i\pi/4} \end{bmatrix}$$

## Multi-Qubit States

A state of more qubits is modelled as a product of all standard bases.

- A two-qubit state is  $\alpha|00\rangle + \beta|01\rangle + \gamma|10\rangle + \theta|11\rangle$ , where  $|\alpha|^2 + |\beta|^2 + |\gamma|^2 + |\theta|^2 = 1$ .
- A three-qubit state is modelled by an eight-element complex vector.

$$|000\rangle = \begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$$|001\rangle = \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

$$|010\rangle = \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \end{bmatrix}$$

...

$$|111\rangle = \begin{bmatrix} 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$$

# Qubits, Gates, Statements

$n \in \mathbb{N}$

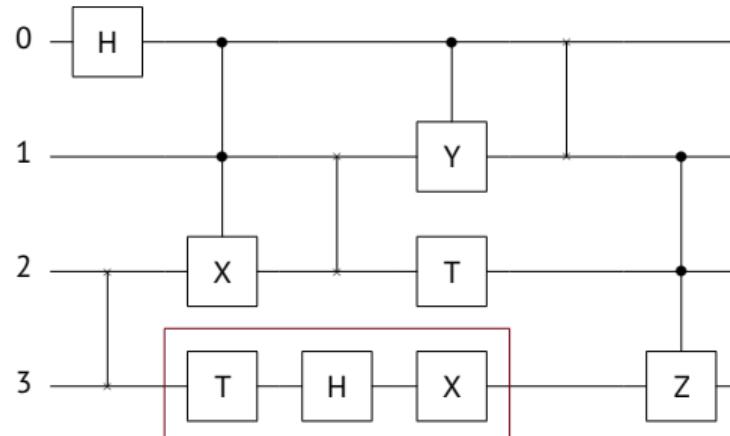
$q \in \text{Qubits} = \mathbb{N}$

$g ::= X | Y | Z | H | T$

$s ::= \text{gate } g \ q \mid \text{swap } q \mid \text{ctrl } n \ g \ q$   
 $\mid s ; s \mid \text{nop}$

Circuits can be represented by statements (different statements may represent the same circuit).

**Notice:** Gates T, H, and X on qubit 3 are subject to *gate fusion!*



(Example circuit)

$s_0 = \text{gate H } 0 ; \text{swap } 2 ;$   
 $\text{ctrl } 2 \ X \ 0 ; \text{gate T } 3 ;$   
 $\text{swap } 1 ; \text{gate H } 3 ;$   
 $\text{ctrl } 1 \ Y \ 0 ; \text{gate T } 2 ; \text{gate X } 3 ;$   
 $\text{swap } 0 ; \text{ctrl } 2 \ Z \ 1$

## Multi-Qubit Circuits

Multi-qubit circuits are modelled using **matrix multiplication** (horizontal composition) and **tensor products** (vertical composition).

$$A \otimes B = \begin{bmatrix} a_{11}B & \cdots & a_{1n}B \\ \vdots & \ddots & \vdots \\ a_{m1}B & \cdots & a_{mn}B \end{bmatrix} \quad A \otimes \mathbf{1} = A = \mathbf{1} \otimes A$$

Special matrices are used for specifying **qubit swapping** ( $SW$ ) and for specifying so-called **control gates** ( $C_n U$ ).

$$SW = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \quad I_n = \begin{cases} \mathbf{1} & \text{if } n = 0 \\ I \otimes I_{n-1} & \text{otherwise} \end{cases} \quad C_0 U = U$$
$$C_n U = \begin{bmatrix} I_n & 0 \\ 0 & C_{n-1} U \end{bmatrix}$$

## Semantics (denotation)

Single-qubit gates  $g$  are denoted, written  $\llbracket g \rrbracket$ , by their corresponding  $2 \times 2$  unitary matrix:

$$\llbracket I \rrbracket = I \quad \llbracket X \rrbracket = X \quad \llbracket Y \rrbracket = Y \quad \llbracket Z \rrbracket = Z \quad \dots$$

A  $k$ -qubit circuit (i.e., a statement)  $s$  is denoted by a unitary matrix  $\llbracket s \rrbracket_k : \mathbb{C}^{2^k \times 2^k}$ :

$$\begin{aligned}\llbracket \text{gate } g \text{ } q \rrbracket_k &= I_q \otimes \llbracket g \rrbracket \otimes I_{k-q-1} \\ \llbracket \text{swap } q \rrbracket_k &= I_q \otimes SW \otimes I_{k-q-2} \\ \llbracket \text{ctrl } n \text{ } g \text{ } q \rrbracket_k &= I_q \otimes C_n \llbracket g \rrbracket \otimes I_{k-q-n-1} \\ \llbracket s_1; s_2 \rrbracket_k &= \llbracket s_2 \rrbracket_k \llbracket s_1 \rrbracket_k \\ \llbracket \text{nop} \rrbracket_k &= I_k\end{aligned}$$

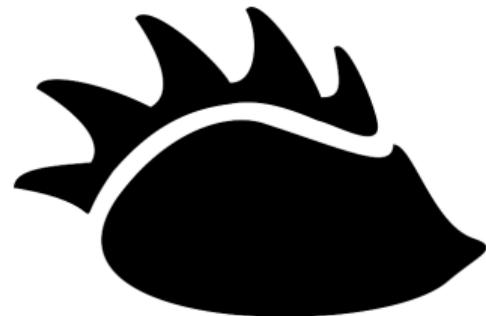
For evaluating a  $k$ -qubit circuit  $s$  on a state vector  $v : \mathbb{C}^{2^k}$ , we can compute  $\llbracket s \rrbracket_k v$ .

Can we do better?

## Futhark: A Data-Parallel Array Language

Futhark is a data-parallel array language featuring standard operations on arrays:

map	:	$\forall n\alpha\beta.(\alpha \rightarrow \beta) \rightarrow [n]\alpha \rightarrow [n]\beta$
transpose	:	$\forall mn\alpha.[m][n]\alpha \rightarrow [n][m]\alpha$
flatten	:	$\forall mn\alpha.[m][n]\alpha \rightarrow [m * n]\alpha$
unflatten	:	$\forall mn\alpha.[m * n]\alpha \rightarrow [m][n]\alpha$
vec	:	$\forall mn\alpha.[m][n]\alpha \rightarrow [n * m]\alpha$
	=	$\text{flatten} \circ \text{transpose}$
unvec	:	$\forall mn\alpha.[m * n]\alpha \rightarrow [n][m]\alpha$
	=	$\text{transpose} \circ \text{unflatten}$



<https://futhark-lang.org>

Futhark targets CUDA, OpenCL, C, and Multi-cores.

## Kronecker-Free, Data-Parallel Interpretation

We can *derive* an efficient interpreter  $\mathcal{I}[s]_k v$  inductively:

For any statement  $s$ , integer  $k > 0$ , and state vector  $v : \mathbb{C}^{2^k}$ , we have  $\mathcal{I}[s]_k v = \llbracket s \rrbracket_k v$

The derivation is based on *vectorisation* and properties of matrix multiplication:

$$(A \otimes B) v = \text{vec}(B(\text{unvec } v)A^T) \quad (1)$$

$$AB = (\text{map } A B^T)^T \quad (2)$$

$$(AB)^T = B^T A^T \quad (3)$$

Calculations yield:

$$\begin{aligned}\mathcal{I}[\text{gate } g q]_k v &= ((I_q \otimes g) \otimes I_{k-q-1}) v \\ &\stackrel{1,3}{=} \text{vec}(((I_q \otimes g)(\text{unvec } v)^T)^T) \\ &\stackrel{2}{=} \text{vec}(\text{map } (I_q \otimes g) (\text{unvec } v)) \\ &= \text{let } F u = \text{flatten}(\text{map } \llbracket g \rrbracket (\text{unflatten } u)) \\ &\quad \text{in } \text{vec}(\text{map } F (\text{unvec } v))\end{aligned}$$

Similar calculations for  
swap and ctrl

## Fusion

Futhark fuses succinct gate operations that target identical qubits:

$$\mathcal{I}[\text{gate } g \text{ } q; \text{gate } g' \text{ } q]_k = \mathcal{I}[\text{gate } g'' \text{ } q]_k \quad (\llbracket g' \rrbracket \llbracket g \rrbracket = \llbracket g'' \rrbracket \wedge k > q)$$

The above identity can be derived from the definition of  $\mathcal{I}[\cdot]_k$ , properties of data-parallel operations, and from *map fusion*:

$$\text{map } f \circ \text{map } g = \text{map } (f \circ g)$$

Other fusion identities hold for consecutive `cntrl` gates performed on identical qubits and for consecutive `swap` gates.

# The dq Futhark Library

The **dq** Futhark library provides rich functionality for programming circuits.

```
module type gates = {
  type c           -- complex numbers
  type q = i64     -- qubits
  type st[n] = [2**n]c -- state vectors
  type^ stT[n] = *st[n] → *st[n]
  ...
  val gateH [n] : q → stT[n]
  val cntrlX[n] : (m:i64) → q → stT[n]
  val swap [n] : q → stT[n]
  val swap2 [n] : (q:q) → (r:q) → stT[n]
  ...
  type ket[n] = [n]i64 -- ket vectors
  val fromKet[n] : ket[n] → *st[n]
  val toKet : (n:i64) → (i:i64) → ket[n]
  type dist[n] = [2**n](ket[n],f64)
  val dist [n] : st[n] → dist[n]
  val distmax [n] : dist[n] → (ket[n],f64)
}
```

```
-- Some utility functions
val >*> 'a : (q → *a → *a) → (q → *a → *a)
          → (q → *a → *a)
val |*>   'a 'b : *a → (*a → *b) → *b
val >* 'a 'b 'c : (*a → *b) → (*b → *c)
          → (*a → *c)
val repeat  'a : i64 → (i64 → *a → *a)
          → *a → *a
```

## Notice:

- Stars (\*) on types indicate uniqueness...
- **dq** can serve as a target language or as a source language for directly expressing quantum algorithms.
- <https://github.com/diku-dk/dq>

## Grover's Search Algorithm in Futhark

```
-- Grover's algorithm: search for the
-- index where oracle returns 1, which
-- it does for the binary encoding of
-- the argument (< 2**n) using a sub-
-- linear number of quantum operations.

import "dqfut"
open mk_gates(f64)

def grover_diff [n] : stT[n] =
    repeat n (gateH >*> gateX)
    >* gateH (n-1)
    >* cntrlX (n-1) 0
    >* gateH (n-1)
    >* repeat n (gateX >*> gateH)

def encNum [n] (i:i64) : stT[n] =
    λs → (loop (s,i) = (s,i) for n in n..>0 do
        if i % 2 == 0
        then (gateX (n-1) s, i/2)
        else (s,i/2)
    ).0
```

```
def oracle [n] i : stT[n] =
    encNum i >* cntrlZ (n-1) 0 >* encNum i

def grover (n:i64) (i:i64) : (ket[n], f64) =
    let k = 2**n |> f64.i64 |> f64.sqrt
        |> (*(f64.pi/4)) |> f64.ceil
        |> i64.f64
    let s = fromKet (replicate n 0)
        |*> repeat n gateH
        |*> repeat k (λ_ →
            oracle i >*
            grover_diff)
    in dist s |> distmax
```

### Notice:

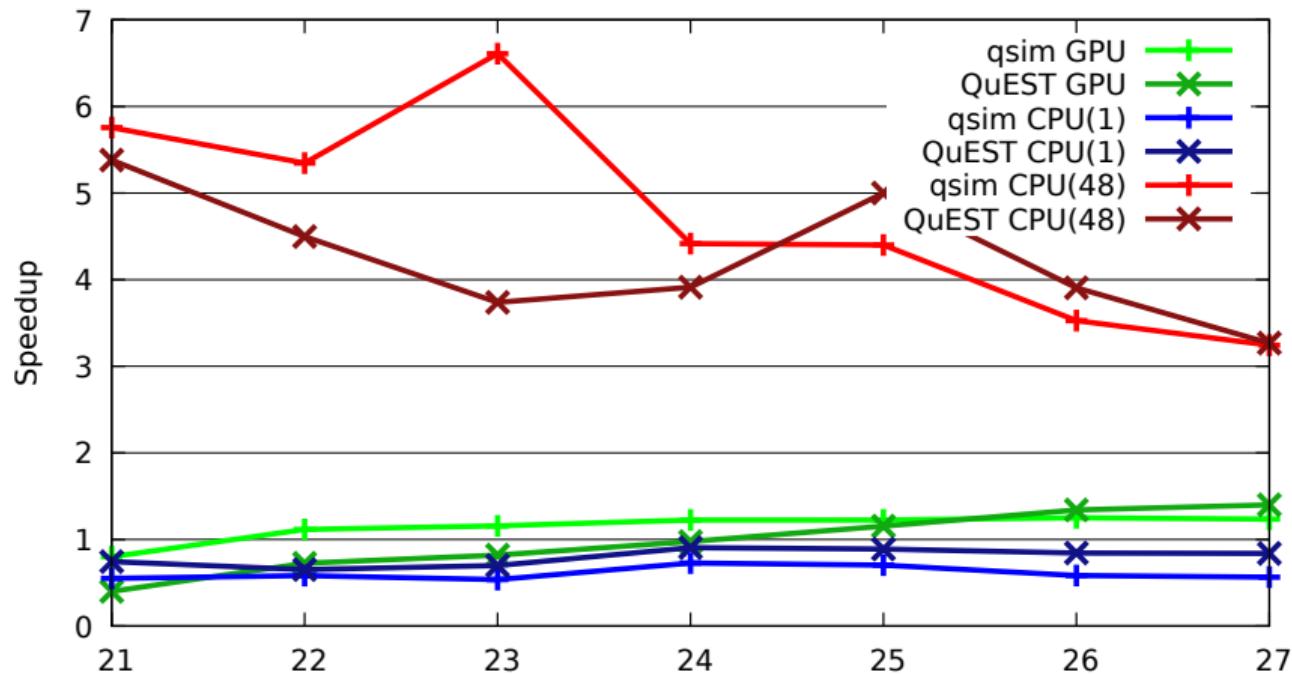
- Boxed gate operations are fused!

We compare with the highly tuned imperative simulators qsim and QuEST for three standard benchmarks (qft, grover, and ghz).

The results are promising:

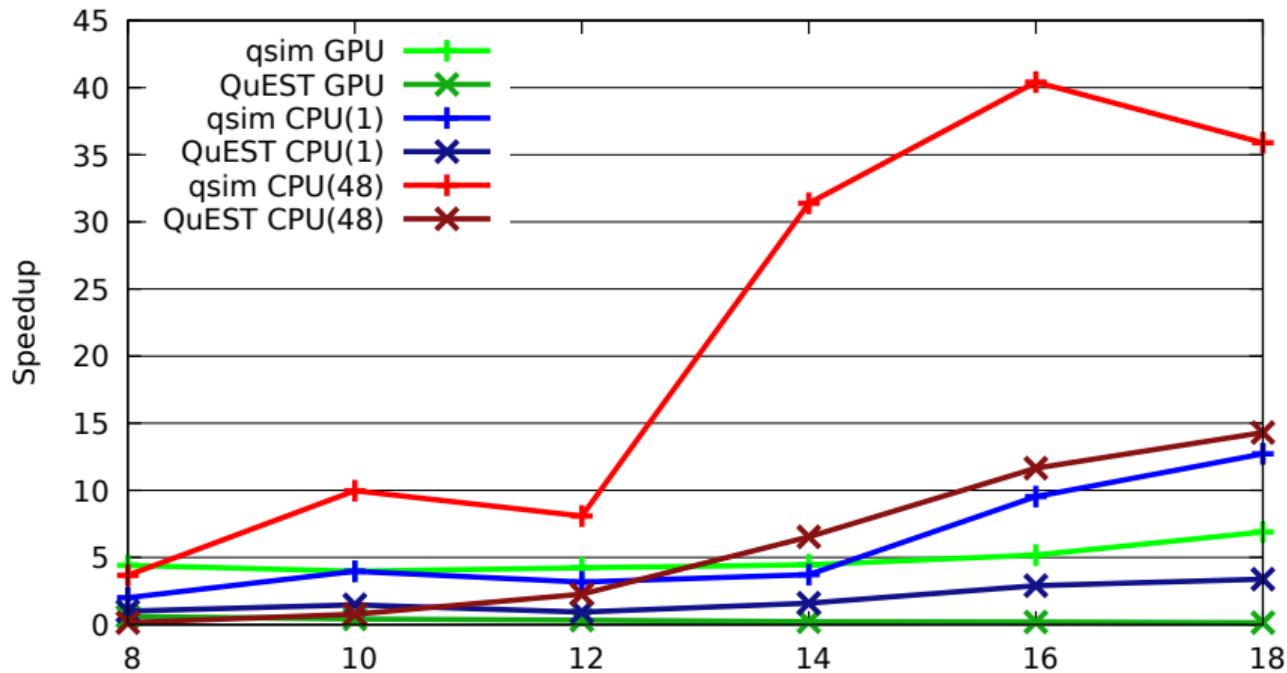
- Fusion leads to significant speedups (i.e., for grover).
- We see slowdowns (sometimes a factor of 10) compared to the highly tuned simulators, which have pointed us towards issues with index transformations in Futhark.

## Benchmark: GHz



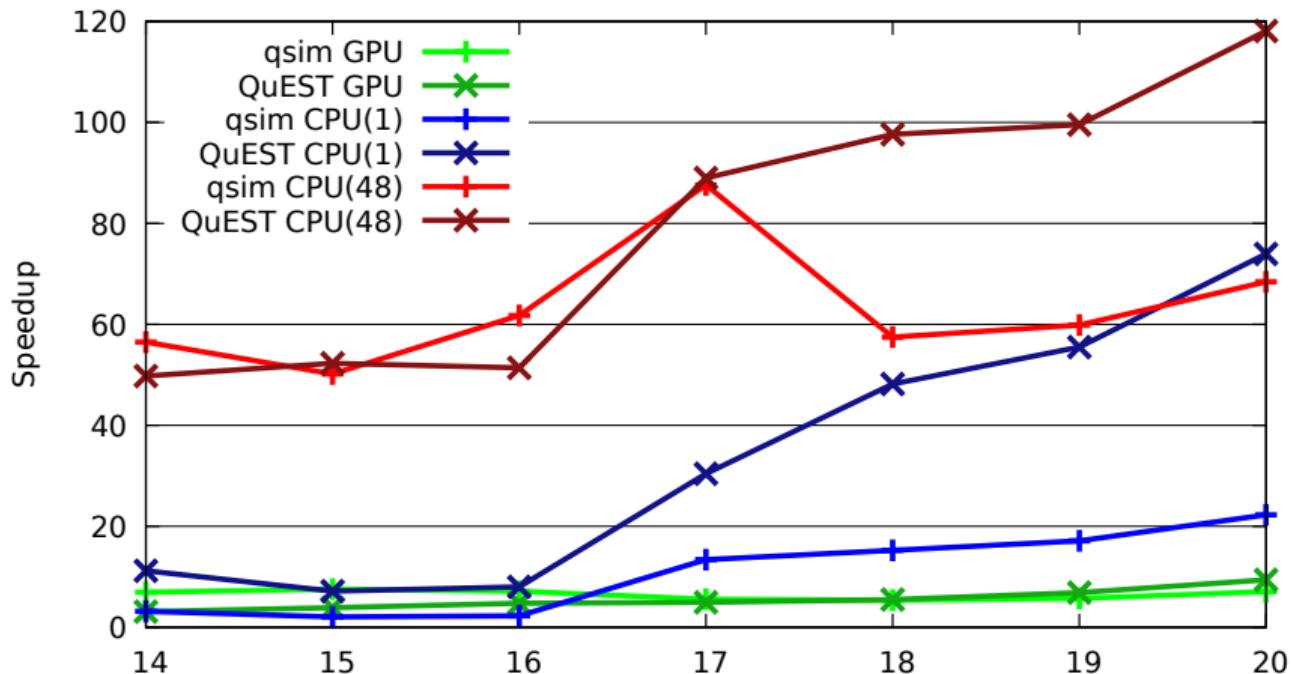
Speedups > 1 indicate that the specified tool performs better than **dq** on the same platform.

## Benchmark: Grover



Speedups > 1 indicate that the specified tool performs better than **dq** on the same platform.

## Benchmark: QFT



Speedups > 1 indicate that the specified tool performs better than **dq** on the same platform.

## More in the Paper

- Extensive derivations of quantum gate operations.
- Derivation of operation for **swapping distant qubits**.
- Derivation showing that gate fusion can be derived from **map fusion**.
- Comparison with **related work** (e.g., other quantum simulation frameworks).
- Directions for **future work** (offline circuit optimisations, lazy state expansion, improvements of Futhark index-optimisations, ...)